

Data Mining For Decision Making

Principal Components Analysis (PCA)

Prof. Varun Dutt

School of Computing and Electrical Engineering
Indian Institute of Technology Mandi, India



Scaling the Heights

Introduction to PCA

Key Ideas:

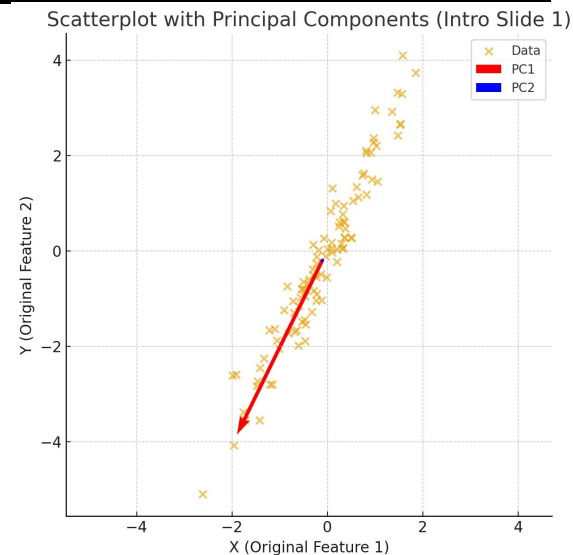
- PCA = a statistical technique for:
 - **Dimensionality reduction** (fewer variables, less redundancy).
 - **Feature extraction** (new meaningful variables = principal components).
- It re-expresses data along new coordinate axes that:
 - Are **uncorrelated** (orthogonal).
 - Capture **maximum variance** in descending order.
- Helps simplify data without losing much information.

Why PCA?

- Many datasets have correlated variables (e.g., height & weight, pixel intensities in images).
- PCA transforms correlated variables → uncorrelated “principal components”.
- First few PCs often capture most of the information.

Applications:

- Face recognition (eigenfaces).
- Image compression.
- Visualization of high-dimensional data.
- Noise reduction.



- Blue points = correlated dataset.
- Red arrow = **PC1 (direction of max variance)**.
- Blue arrow = **PC2 (orthogonal direction, less variance)**.
- This gives a perfect introduction to PCA as “rotating axes to capture maximum spread”.

Variance, Covariance, and Covariance Matrix

Variance (σ^2)

- Measures how spread out a variable is around its mean.
- Formula:

$$Var(X) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$$

- Example: If test scores vary widely, variance is high.

Covariance Matrix

- A square matrix that summarizes variance and covariance of all variables.

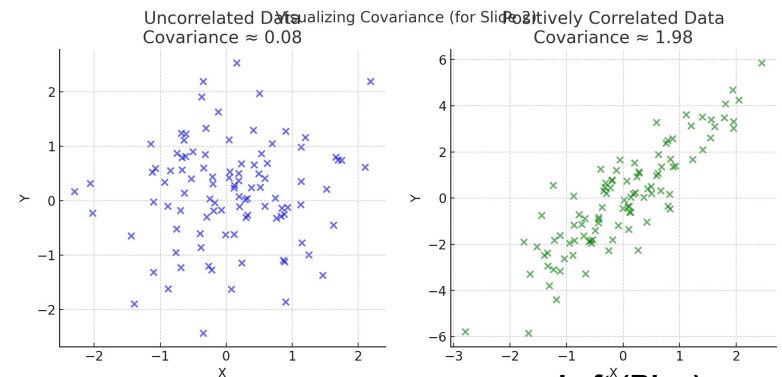
For 2D data:

$$\Sigma = \begin{bmatrix} Var(X) & Cov(X, Y) \\ Cov(X, Y) & Var(Y) \end{bmatrix}$$

- Diagonal entries** = variances.
- Off-diagonal entries** = covariances.

Why Important for PCA?

- PCA rotates the data so that axes (principal components) align with directions of **maximum variance**.
- Covariance matrix is the starting point for finding these directions.



Covariance ($cov(X,Y)$)

- Measures how two variables change together.
- Formula:

$$Cov(X, Y) = \frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})$$

- Interpretation:
 - Positive** \rightarrow X and Y increase together.
 - Negative** \rightarrow one increases, other decreases.
 - Zero** \rightarrow no linear relationship.

Left (Blue):

Uncorrelated data \rightarrow covariance ≈ 0 .

Right (Green):

Positively correlated data \rightarrow covariance > 0 .

PCA Workflow (Step-by-Step)

Step 1: Collect the Data

- Organize dataset into an $n \times d$ matrix (n = samples, d = features).
- Example: 10 points in 2D $\rightarrow 10 \times 2$ matrix.

Step 2: Mean Center the Data

- Subtract the mean of each feature \rightarrow new data with mean = 0.
- Ensures PCA is not biased by absolute position.

Step 6: Choose k Components

- Keep top k eigenvectors \rightarrow reduce dimensionality.
- Balance between dimensionality reduction and information loss.

Step 3: Compute Covariance Matrix

- Build covariance matrix Σ to capture relationships between features.
- Σ dimension = $d \times d$.

Step 4: Find Eigenvalues & Eigenvectors

- Solve equation:

$$\Sigma v = \lambda v$$

- Eigenvectors (v) = directions of new axes (principal components).
- Eigenvalues (λ) = variance captured along each axis.

Step 5: Sort by Variance

- Rank eigenvectors by descending eigenvalues.
- First component = PC1 (maximum variance).
- Second component = PC2, and so on.

Step 7: Transform the Data

- New data = $FeatureVector^T \times MeanAdjustedData^T$.
- Original data \rightarrow expressed in terms of principal components.

Step 1: Example Dataset & Mean Centering

Original Dataset (10 points in 2D)

x	y
2.5	2.4
0.5	0.7
2.2	2.9
1.9	2.2
3.1	3.0
2.3	2.7
2.0	1.6
1.0	1.1
1.5	1.6
1.1	0.9

Step 1: Compute the Mean

- Mean of x:

$$\bar{x} = 1.81$$

- Mean of y:

$$\bar{y} = 1.91$$

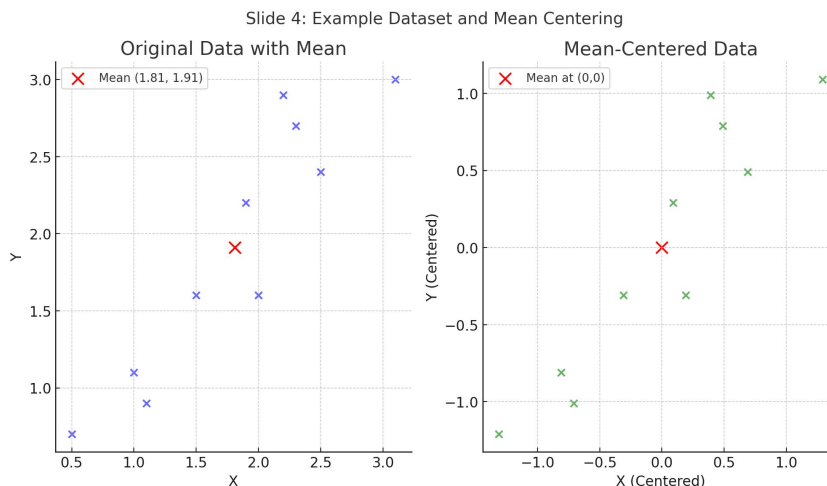
Step 2: Mean-Center the Data

Subtract the mean from each value.

- Example: $(2.5, 2.4) \rightarrow (0.69, 0.49)$.
- After transformation, dataset has **mean = (0,0)**.

Why Mean Centering?

- Moves the dataset to the origin.
- Ensures PCA finds directions of maximum variance independent of absolute location.



Left: Original dataset with mean point marked (red X at (1.81, 1.91)).

Right: Mean-centered dataset (shifted so mean is at the origin).

Step 2: Covariance Matrix

Definition

For 2D data:

$$\Sigma = \begin{bmatrix} \text{Var}(X) & \text{Cov}(X, Y) \\ \text{Cov}(X, Y) & \text{Var}(Y) \end{bmatrix}$$

Where:

- $\text{Var}(X)$ = variance of feature X
- $\text{Var}(Y)$ = variance of feature Y
- $\text{Cov}(X, Y)$ = how X and Y vary together

Computation (from handout dataset)

Using mean-centered data:

- $\text{Var}(X) = 0.6166$
- $\text{Var}(Y) = 0.7166$
- $\text{Cov}(X, Y) = 0.6154$

Thus:

$$\Sigma = \begin{bmatrix} 0.6166 & 0.6154 \\ 0.6154 & 0.7166 \end{bmatrix}$$

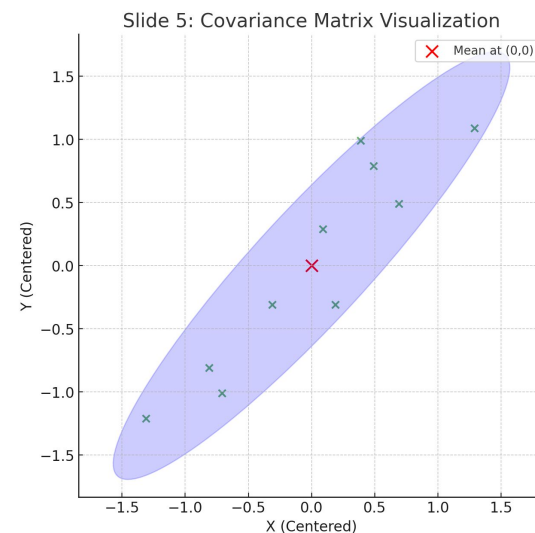
Interpretation

- Positive covariance \rightarrow X and Y increase together.
- Diagonal entries (variances) show spread of each feature.
- Off-diagonal entry shows correlation strength.

Green points =
mean-centered
dataset.

Red X = mean at
(0,0).

Blue ellipse =
covariance ellipse
(spread and
correlation).



Step 3: Eigenvalues & Eigenvectors

Eigenvalue Equation

We solve:

$$\Sigma v = \lambda v$$

For covariance matrix:

$$\Sigma = \begin{bmatrix} 0.6166 & 0.6154 \\ 0.6154 & 0.7166 \end{bmatrix}$$

Step 1: Characteristic Polynomial

$$\det(\Sigma - \lambda I) = 0$$

$$\det \begin{bmatrix} 0.6166 - \lambda & 0.6154 \\ 0.6154 & 0.7166 - \lambda \end{bmatrix} = 0$$

$$(0.6166 - \lambda)(0.7166 - \lambda) - (0.6154)^2 = 0$$

Step 2: Expand

$$\lambda^2 - (0.6166 + 0.7166)\lambda + (0.6166)(0.7166) - (0.6154)^2 = 0$$

$$\lambda^2 - 1.3332\lambda + 0.0630 = 0$$

Step 3: Solve Quadratic

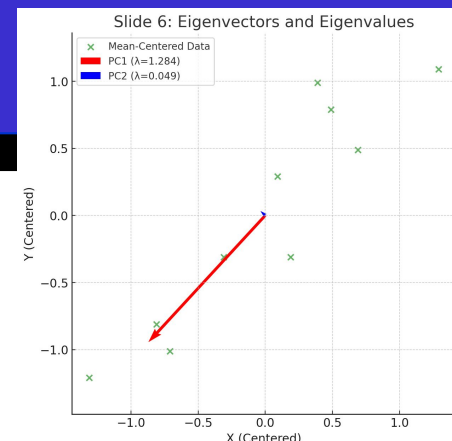
$$\lambda = \frac{1.3332 \pm \sqrt{1.3332^2 - 4(0.0630)}}{2}$$

$$\lambda_1 = 1.284, \lambda_2 = 0.049$$

Green points = mean-centered dataset.

Red arrow (PC1) = direction of maximum variance ($\lambda_1 = 1.284$).

Blue arrow (PC2) = orthogonal minor component ($\lambda_2 = 0.049$).



Step 4: Find Eigenvectors

For $\lambda_1 = 1.284$:

$$\begin{bmatrix} -0.6674 & 0.6154 \\ 0.6154 & -0.5674 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} = 0$$

Solution: $y = 1.085x \rightarrow$ eigenvector $v_1 = (-0.678, -0.735)$.

For $\lambda_2 = 0.049$:

Solution: $y = -0.923x \rightarrow$ eigenvector $v_2 = (-0.735, 0.678)$.

Interpretation

- **PC1 ($\lambda_1=1.284, v_1$):** direction of maximum variance (~96%).
- **PC2 ($\lambda_2=0.049, v_2$):** orthogonal, minor variance (~4%).

Step 4: Projection onto Principal Components

Why Project?

- We want to express the original data in terms of the new axes (principal components).
- This gives a rotated coordinate system where:
 - PC1 = maximum variance direction.
 - PC2 = orthogonal direction with minimal variance.

Transformation Equation

If X = mean-centered data matrix,
and W = matrix of eigenvectors (columns),

then transformed data:

$$Y = XW$$

- Each row of Y = new coordinates of a data point in PC space.
- If we keep only PC1 \rightarrow 1D representation (dimensionality reduction).

Example (Handout Data)

Original (mean-centered point):

(0.69, 0.49)

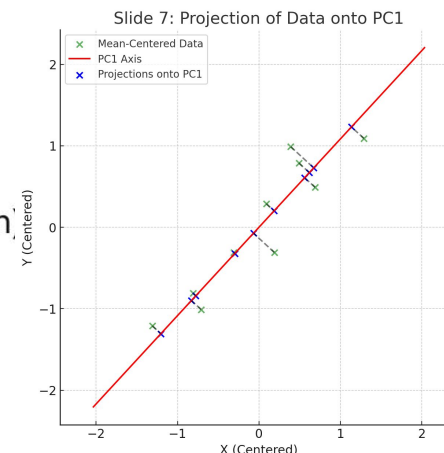
Dot product with PC1 $(-0.678, -0.735)$:

$$0.69(-0.678) + 0.49(-0.735) = -0.827$$

So in PC1 space, the point projects to ≈ -0.827 .

Interpretation

- Data "shadow" onto PC1 captures **most variance (~96%)**.
- Using only PC1: reduces 2D \rightarrow 1D while retaining main pattern.
- Using both PCs: exact reconstruction.



Green points: Original mean-centered data.

Red line: PC1 axis (direction of maximum variance).

Blue points: Projections of data onto PC1.

Dashed lines: Show how each original point is projected down to PC1.

Step 5: Variance Explained by PCs

Eigenvalues Recap

- Each eigenvalue λ = variance captured along its eigenvector (PC).
- Larger $\lambda \rightarrow$ more important component.

From our dataset:

- $\lambda_1 = 1.284$ (PC1)
- $\lambda_2 = 0.049$ (PC2)

Explained Variance Ratio

$$\text{Explained Variance Ratio} = \frac{\lambda_i}{\sum \lambda}$$

- PC1:

$$\frac{1.284}{1.284 + 0.049} = 0.963 \text{ (96.3\%)}$$

- PC2:

$$\frac{0.049}{1.284 + 0.049} = 0.037 \text{ (3.7\%)}$$

Interpretation

- PC1 captures almost all of the variance ($\approx 96\%$).
- PC2 adds very little information ($\approx 4\%$).
- So projecting onto PC1 alone still preserves most structure in data.

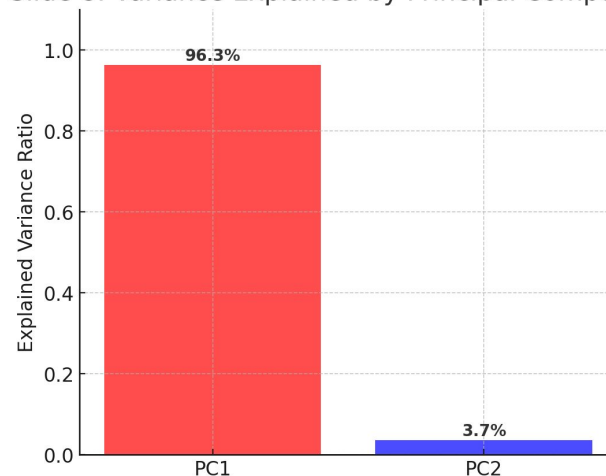
Applications

Dimensionality Reduction:

- Keep only PCs with large eigenvalues.
- Discard PCs with tiny eigenvalues (mostly noise).

- **Data Compression:** Store fewer dimensions with minimal info loss.

Slide 8: Variance Explained by Principal Components



PC1 (red): explains **96.3%** of the variance.

PC2 (blue): explains only **3.7%**.

Step 6: Transforming Data to PC Space

Transformation Equation

$$Y = XW$$

- X : mean-centered dataset
- W : eigenvector matrix (columns = eigenvectors)
- Y : transformed dataset in terms of principal components

Example (Handout Data)

For point (0.69, 0.49):

- Projection onto PC1: ≈ -0.827
- Projection onto PC2: ≈ -0.175

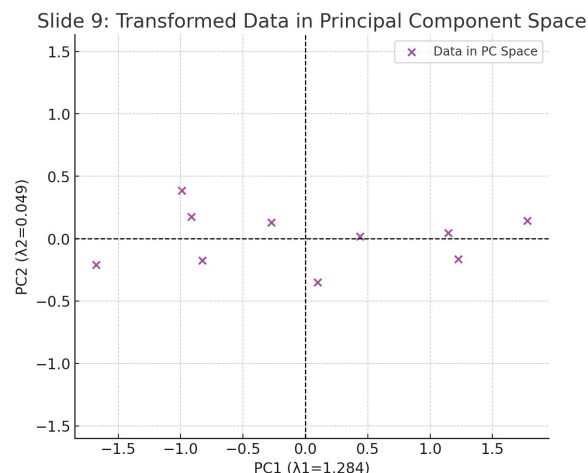
So in PC space: $(-0.827, -0.175)$.

PC Space

- New coordinates = linear combinations of original features.
- Axes are now **uncorrelated**.
- First axis (PC1) shows most variance.

Interpretation

- Scatterplot of data in PC1–PC2 space looks “uncorrelated.”
- PC1 axis is stretched \rightarrow dominant.
- PC2 axis is compressed \rightarrow small variance.



- Scatterplot of the dataset in **PC space (PC1 vs PC2)**.
- Most variance is clearly along the **PC1 axis**, while **PC2 shows very little spread**.
- Confirms that dimensionality reduction to PC1 alone is effective.

Step 7: Reconstruction from PCs

Full Reconstruction (PC1 + PC2)

- If we use **all PCs**, the transformation is invertible:

$$X \approx YW^T + \text{mean}$$

- This gives back the **original dataset exactly**.

Reduced Reconstruction (PC1 only)

- Keep only PC1 → project data to 1D, then map back:

$$X_{approx} \approx Y_{PC1} \cdot W_{PC1}^T + \text{mean}$$

- Data lies along the PC1 line.

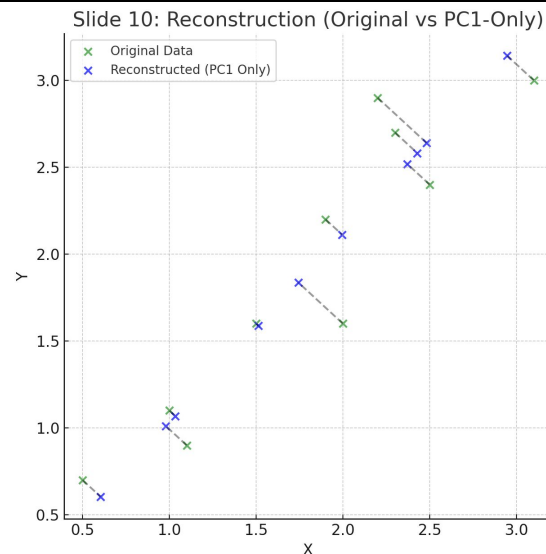
Example

Point (0.69, 0.49):

- Projection on PC1: -0.827.
- Reconstructed using PC1 only: $\approx (0.57, 0.62)$.
- Close to original, but slightly shifted toward PC1 line.

Interpretation

- Using **all PCs**: no information loss.
- Using **PC1 only**: big compression, small loss.
- For this dataset: PC1 already captures **96% of the variance** → little information lost.



Green points: Original dataset.

Blue points: Reconstruction using **PC1 only**.

Dashed lines: Show the small error (distance lost along PC2).

Applications of PCA

1. Data Compression

- Reduce dimensionality while keeping most information.
- Example:
 - Images → keep only top PCs ("eigenimages").
 - Store fewer numbers, yet image is still recognizable.

2. Noise Reduction

- Small eigenvalues often capture random noise.
- By discarding them, PCA denoises data.
- Example: ECG or EEG signals — remove noisy components.

3. Pattern Recognition

- PCA reveals underlying structure.
- **Face Recognition (Eigenfaces):**
 - Each face = combination of principal components (eigenfaces).
 - Compare faces in reduced PC space → faster, robust recognition.


4. Data Visualization

- High-dimensional data → reduce to 2D or 3D.
- Example: visualizing gene expression profiles or word embeddings.

Key Takeaway

- PCA is not just math → it's a powerful tool for:
 - **Simplification**
 - **Compression**
 - **Visualization**
 - **Recognition**

Slide 11: PCA for Image Compression



Left: Original 8×8 digit image.

Next: Reconstructions with only **2 PCs**, **4 PCs**, and **8 PCs**.

Shows how fewer principal components still capture the main structure, but with some loss in detail.

PCA: Summary & Key Points

Key Insights

- **Principal Components (PCs):**
 - New orthogonal axes capturing max variance.
- **Eigenvalues:** tell how much variance each PC explains.
- **Explained Variance Ratio:** helps decide how many PCs to keep.
- **Dimensionality Reduction:** keep only top PCs → smaller, simpler dataset.

Takeaway Message

- PCA = **rotate and compress** data.
- Captures essential structure while reducing complexity.
- Widely used in **ML, data science, and pattern recognition**.